

# OPI API Meeting

Billy McFall, Donald Hunter and Maryam Tahhan

March 23, 2023



# Vendor Dataplane View

OPI API used to define intent

- Controls both kernel networking and vendor dataplane
- Support for routing & flow based intent

Hardware abstraction API used to mirror state between kernel and vendor dataplane

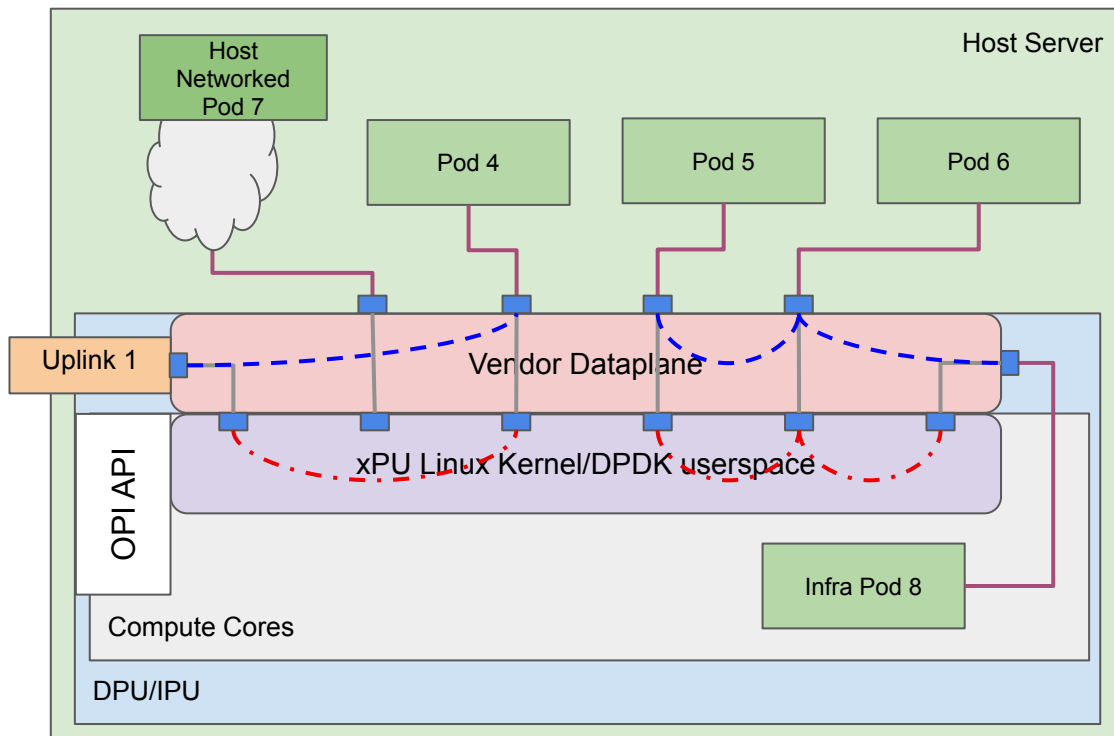
- netlink / tc-flower / rte-flow

Control plane traffic handled by kernel

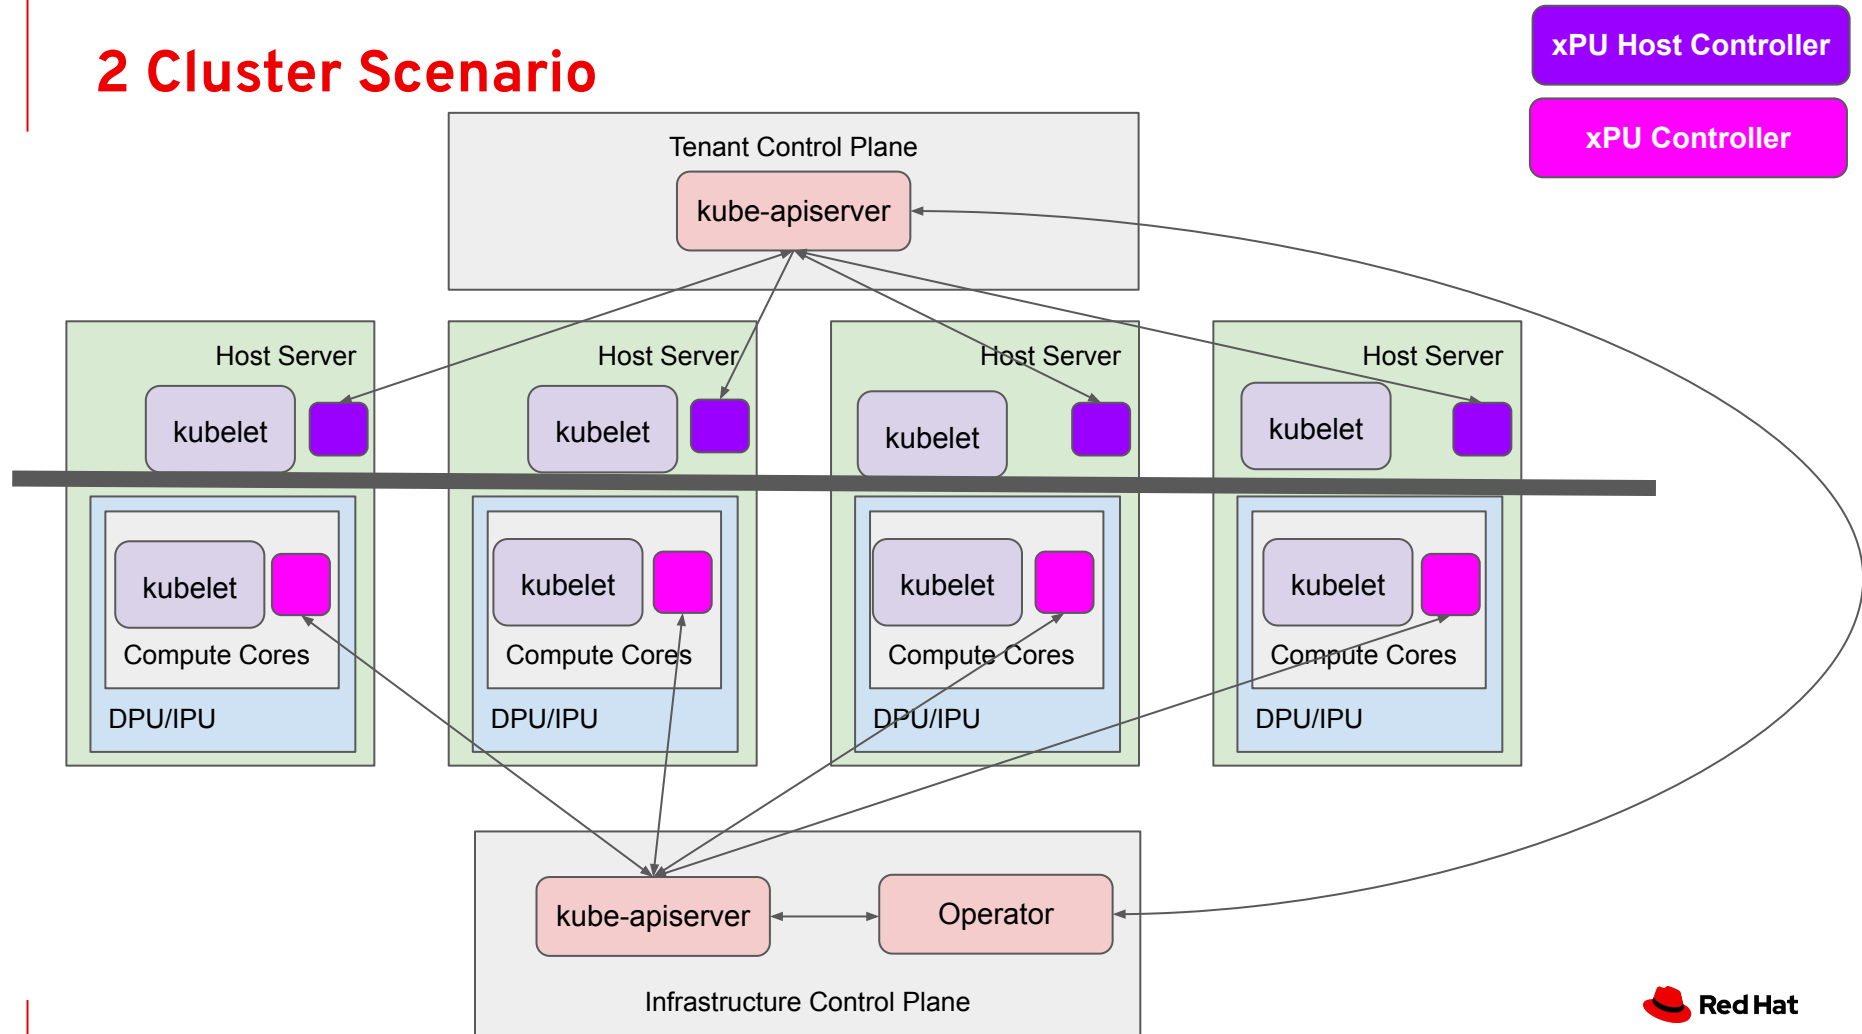
Fast path in vendor dataplane - - - - .

Fallback to slow path in kernel - . - . - .

## IPU Deployment



## 2 Cluster Scenario



# Entities xPU

## Operator

- ▶ Deploys the xPU Agent Daemonsets
- ▶ CRUDs OPI CRDs.
- ▶ Monitors Kubernetes resources in the tenant cluster.
- ▶ Translates incoming resources into OPI CRDs.
- ▶ Syncs OPI CRDs to the Tenant Cluster.

## xPU Agent Daemonset

- ▶ Monitors OPI CRDs.
- ▶ Translates those CRDs into OPI API calls (which it also invokes).
- ▶ CRUDs OPI CRDs.

# Entities Host

## Device Plugin

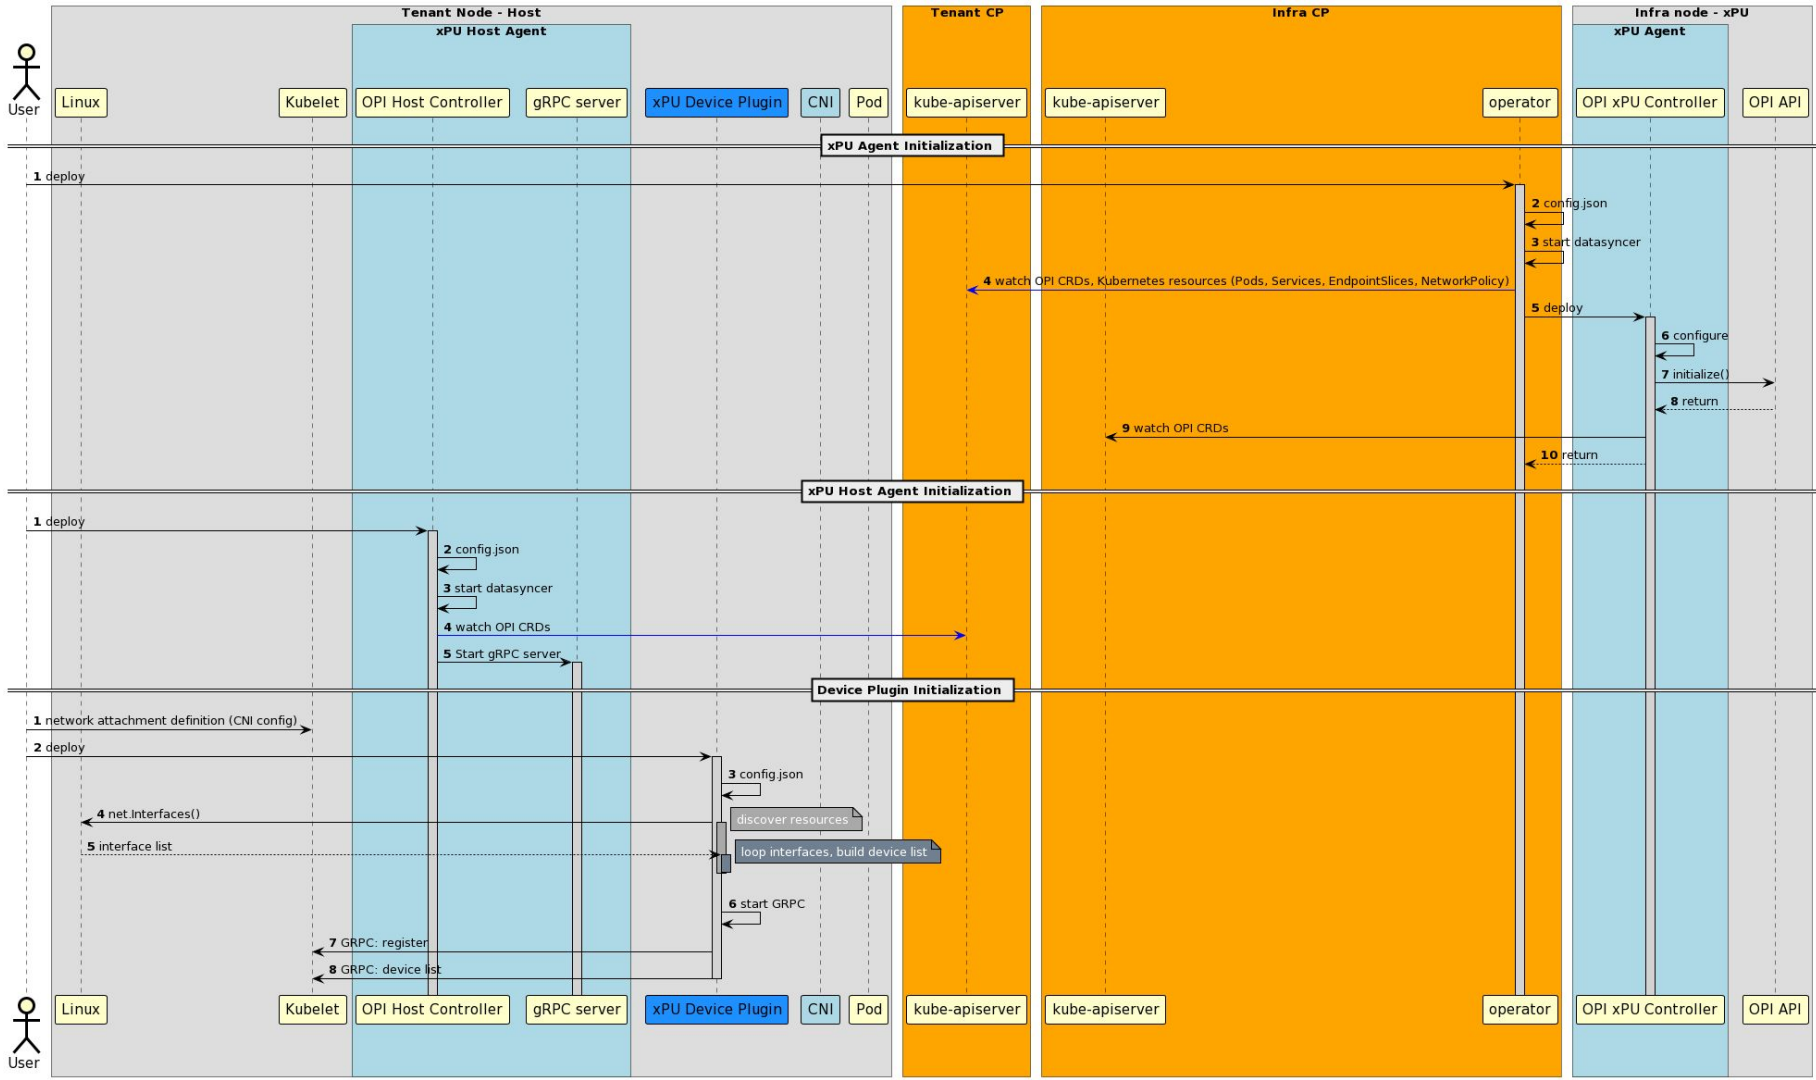
- ▶ Provisions and advertises the VFs/SFs to Kubernetes.
- ▶ Interacts with the xPU Host agent to invoke the creation of OPI netdev CRDs on Pod allocation.

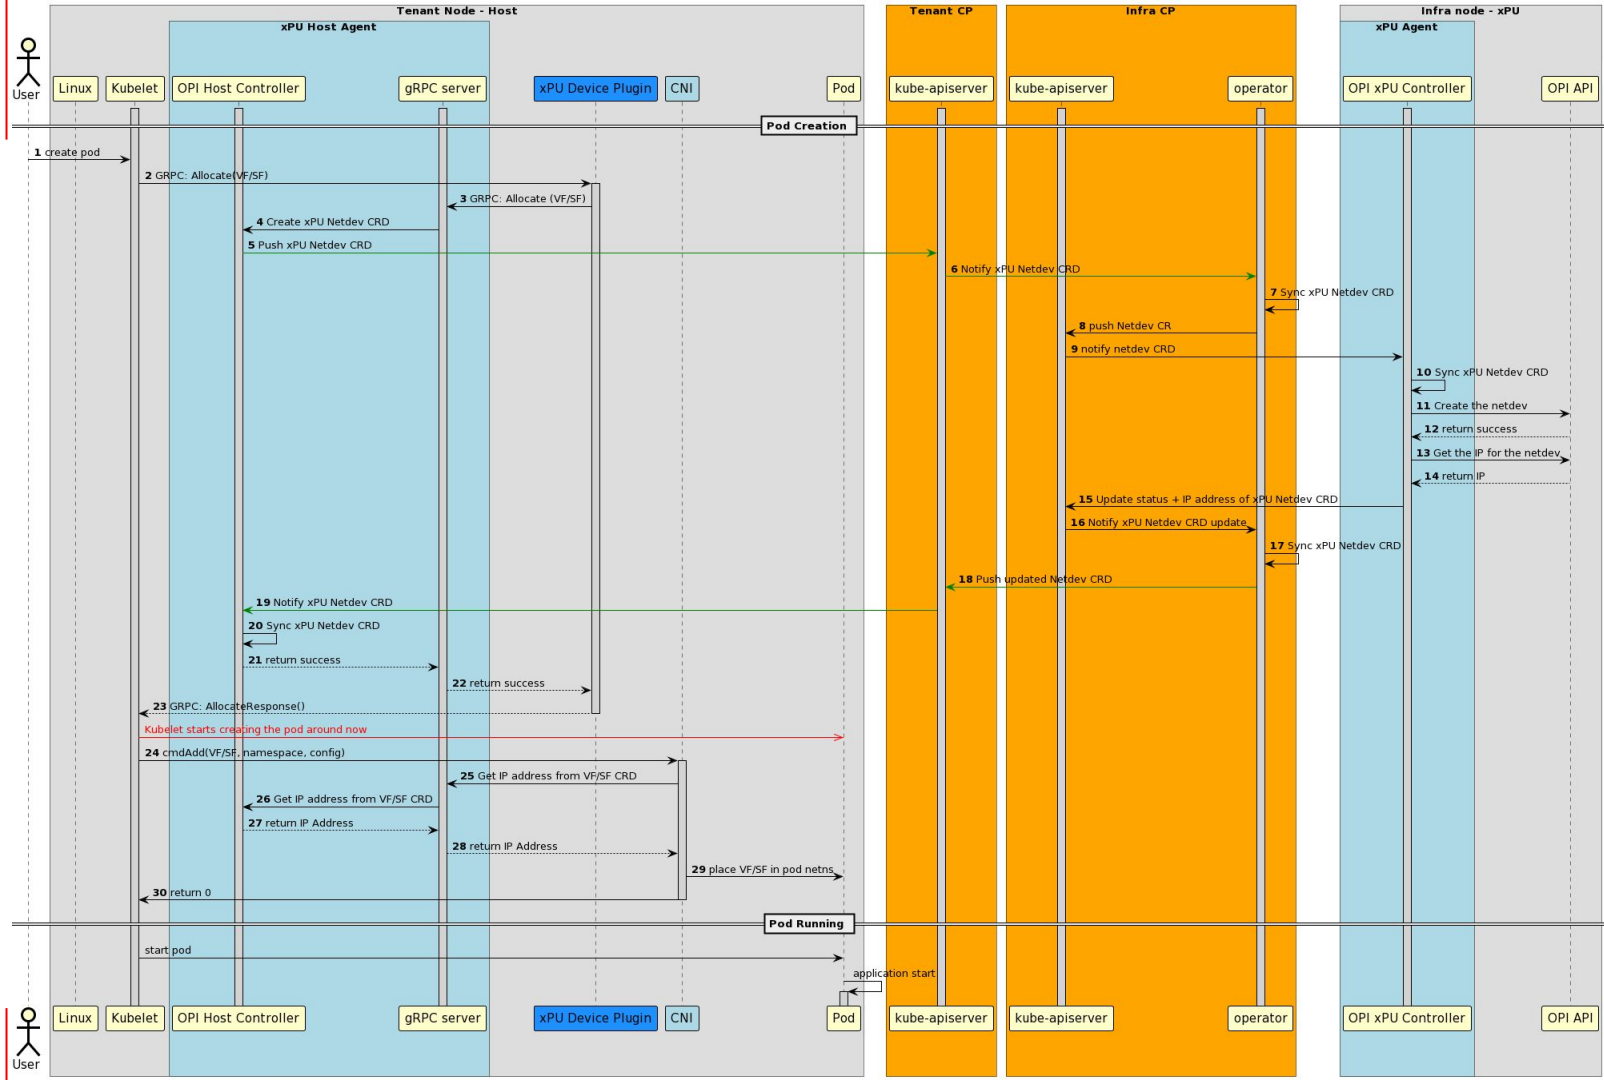
## CNI

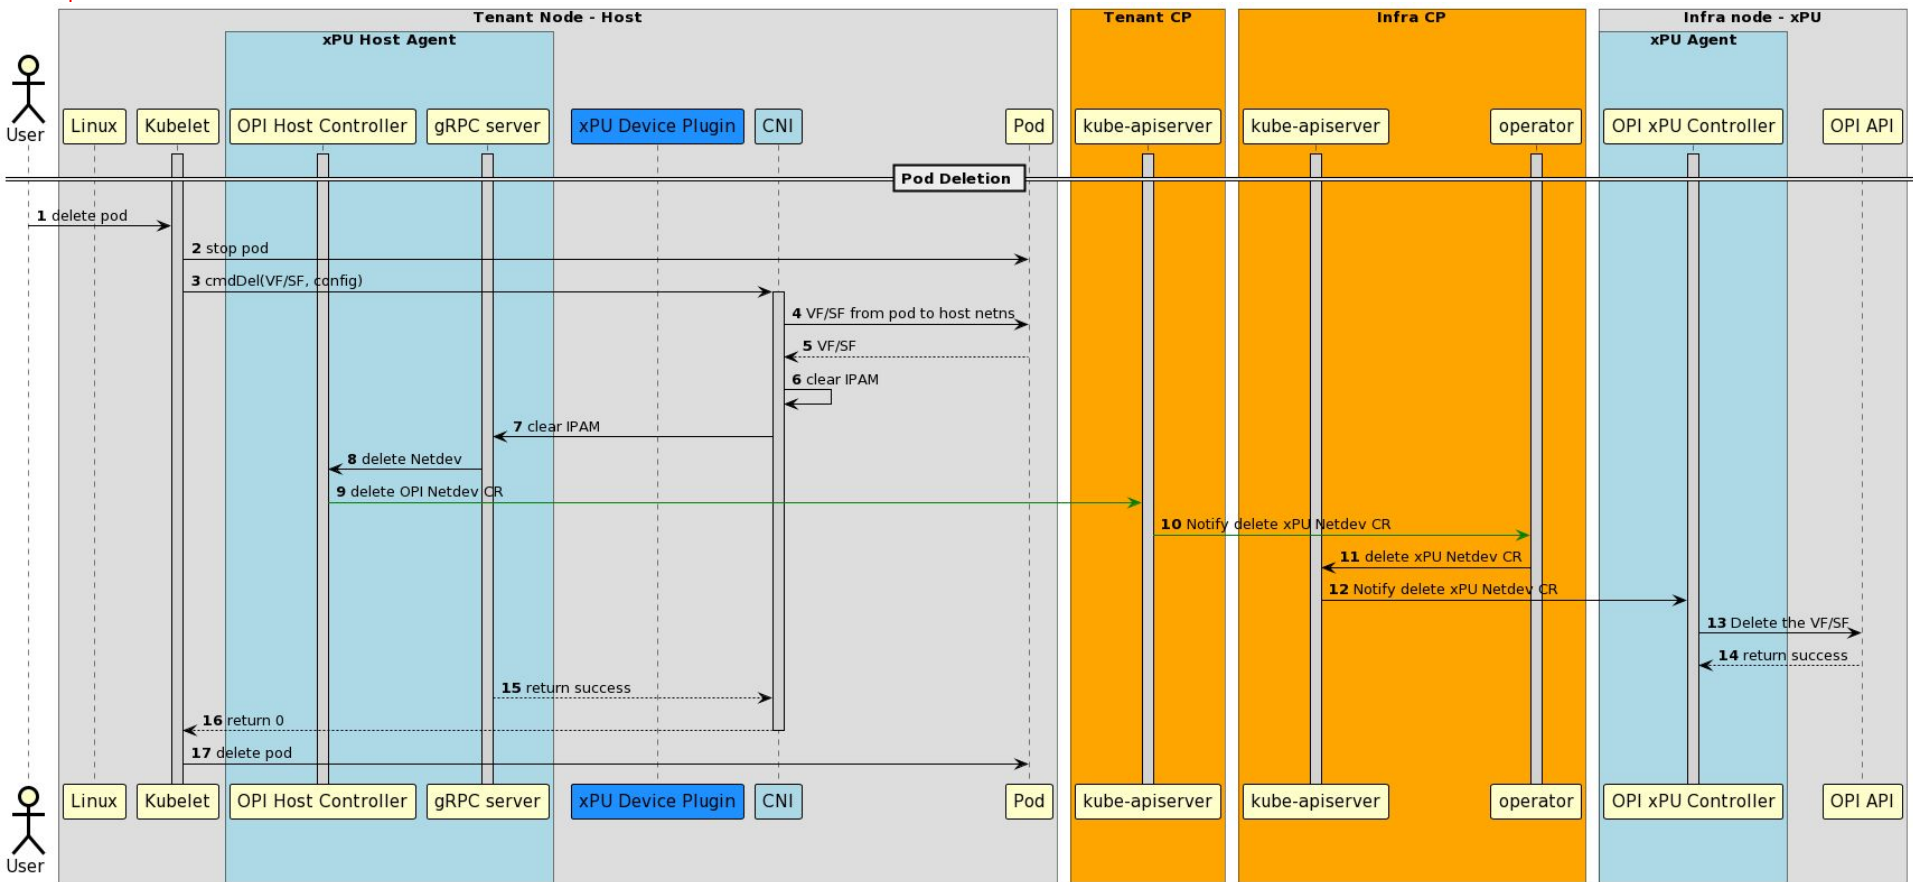
- ▶ Configures the interface with the allocated IP address
- ▶ Moves the VF/SF from the Host network namespace to the Pod Network namespace (vice versa).

## xPU Host Agent Daemonset

- ▶ CRUDs OPI CRDs (from the Kube-apiserver or the Broker).
- ▶ Processes requests from the Device Plugin or CNI via a gRPC server.









## Links to sequence diagrams

<https://github.com/redhat-et/opi-k8s-networking/blob/5d3a466cc2ac0ed28c8afe524e3e43fc749967a1/sequence-diagrams/images/host-provisioning-xPU-VF-init.png>

<https://github.com/redhat-et/opi-k8s-networking/blob/5d3a466cc2ac0ed28c8afe524e3e43fc749967a1/sequence-diagrams/images/host-provisioning-xPU-VF-pod-creation.png>

<https://github.com/redhat-et/opi-k8s-networking/blob/5d3a466cc2ac0ed28c8afe524e3e43fc749967a1/sequence-diagrams/images/host-provisioning-xPU-VF-pod-deletion.png>